AUtoMLy | Progress Log 2

INAI 6986 · Week 8: October 5–11, 2026 · J.O.

Model Evaluation week · Updated October 16 · Local syllabus: Log 2 October 11; Milestone 2 October 25. Canvas not re-verified.

Log Number And Period	Log 2 · Week 8 (October 5–11). The main status below covers that period. Later source, recovery and document work appears only in the dated addendum.
Planned Versus Completed	Planned: use completed training to diagnose errors and begin governed evaluation. The 16 training runs were completed September 29–October 1, before this period. During Week 8: Batch A dispositions received October 5–6 (27 responses including 16 duplicate groups); Batch B submitted October 7 at revision 11 (30 cards, 180 ratings); one bounded prompt refinement completed (8 paired questions, 16 outputs), not promoted. Final evaluation was not run.
Data Readiness	1,999 answerable records; historical document-grouped train/validation/test counts 1,684/124/191. All 30 original controls map to historical-test source documents. That test is exposed, not untouched. Human duplicate rulings are complete; application of versioned remedies to data/selected weights is not evidenced. The fresh benchmark was not yet created.
Decisions Made	October 8: retain Nemotron N=1,600 checkpoint-400 as the proposed candidate, not a factual winner; three adapters tied at 4/24 supported outcomes in the advisory development probe. Preserve the original student controls. The later local-AI/source-verified benchmark authorization is recorded in the dated addendum, not backdated into Week 8; no instructor waiver is implied.
Blockers And Risks Triggered	Advisory probe: 192 slots, 150 judgments, 19 generation-budget failures and 23 parser failures; human accuracy not inferred. Advisory labels include 57 wrong values, 32 abstention issues and 26 invented rules. Responses: withhold promotion, retain failed attempts, withdraw biased scorer v1, and reserve new human output adjudication for the governed final run.
Next Period: Three Priorities	Updated October 16 priorities for Week 9 (Oct 12–18): 1) authenticate/exposure-audit fresh sources and independently verify local-AI benchmark keys; 2) prove native matched candidate/base operation and reconcile primary-metric/human-sample protocol; 3) seal exact membership/settings/robustness, then execute the authorized comparison once and obtain new output ratings. Milestone 2 remains incomplete until actual results exist. Thursday instructor check-in: escalate missing final evidence.
Hours Logged	20 personal hours — student estimate supplied October 14; not measured worker runtime or a stopwatch trial. No independently verified allocation solely to October 5–11 is inferred.

AI Use Disclosure

Hermes Agent (Luna draft; GPT-6.1 Sol final review) assembled and corrected this log; local Qwen supplied development advice. I conceived and directed AUtoMLy, authored the original 30 controls, supplied recorded human ratings and the hours estimate, checked evidence and retain final responsibility. Agents did not supply human ratings or final-test results.

AUtoMLy | Log 2 Evidence And Accountability

Week 8 record with separately dated follow-up

J.O.

Dated Follow-Up — October 14–16

October 14: the separately authorized fresh items are local-AI-authored and source-verified, not new independent student-authored controls; no additional student question-writing or instructor waiver is implied. Identical official-source passages plus questions replaced question-only input for the fixed Nemotron400 versus untouched-base final comparison. Conditional advance authority requires reviewed tools, verified source/keys and immutable assets/settings; no additional freeze/run prompt is needed if those exact conditions pass. This does not waive later human output ratings.

October 16, 09:14 CDT: authorized recovery and real-installation integration passed. All 226 copied records were preserved, 225 prior run keys retired and old admissions refused; source-only preflights authenticated 8 originals/28 physical pages twice with zero model calls. The earlier blocked transition was superseded. Recovery is not a completed benchmark or model evaluation.

Current state at 2026-10-16 15:03:32 CDT: no accepted fresh cohort, native matched-model proof, immutable final seal, final outputs, final robustness delta or new human output ratings. Source audit: 11 of 12 documents have no project-relative exposure evidence; one pharmacy source requires review. The current evidence-replay notebook recomputes development results, not a final test.

Evidence And Credit Rationale

This is the combined PM Plan/Logs rubric, not an invented standalone Log 2 rubric. The original row is literally Progress Log 1. Log 2 contributes the reusable status standard; plan-level credit remains anchored to the PM Plan, not this log alone.

Work Plan / CPMAI (25)	Page 1: planned/completed and three priorities; PM Plan phased schedule.
Risk Register (20)	Page 1: fired risks, counts and actions; PM Plan risk likelihood, impact and triggers.
Data / Environment (20)	Page 1: corpus/split and exposure limitations; current executed replay and versioned records.
Cross-Functional / Evaluation (15)	Page 1: analyst/human-review dependency; addendum: conditional run, protocol and readiness gates.
Progress Log 1 (15)	Evidence stays with Log 1. This Log 2 independently follows all seven reusable template fields.
Communication / Format (5)	One-page core; dated addendum; complete disclosure. No backdated training or invented measured hours.

AI Assistance And Human Accountability

Luna and local Qwen assisted the draft; GPT-6.1 Sol personally rechecked receipts, corrected dates/authorship/gates and inspected final layouts. Local workers performed data generation and development scoring. I chose the architecture, models, metrics and scope, configured hardware, directed and troubleshot agents, supplied human dispositions, chose stop/quarantine actions and retain responsibility for the submitted conclusions. New AI-assisted wording is never independent student authorship.

Evidence: evidence_snapshot_20261016.json, SHA-256 f0c66d60b669e1d26804ddc85802a66ea76f79289feb97c2d667a2f4ce790d1f; current_evidence_replay_20261016.ipynb. Repository: https://github.com/porespellar/INAI6986_project_deliverables . Source receipts remain in the private source-of-record repository; instructor access and Canvas submission are not verified.