AUtoMLy | Progress Log 2

INAI 6986 · Week 8: October 5–11, 2026 · J.O.

Model Evaluation week · Log due October 11; updated October 18 · Week 8 facts unchanged; later events carry their dates.

Log Number & Period	Log 2 · Week 8, October 5–11, 2026.
Planned vs. Completed	Planned (PM Plan, Week 8): evaluation metrics, validation strategy and error analysis for the trained models. Completed: (1) Week 8 error review: I rated 30 Nemotron outputs (15 validation questions, base and adapter; 180 ratings), October 6–7. (2) Rulings on all 16 cross-split duplicate-answer groups and one near-duplicate pair, October 5–6. (3) One prompt-refinement experiment on 8 questions, October 7; it failed and was not promoted. (4) A predeclared comparison of all four model families, base and adapter, on the same 24 validation questions (192 outputs), October 8. Not completed: the one-time final test, moved to Weeks 9–10 because the fresh benchmark is not built yet.
Data-Readiness Status	1,999 answerable records; split by document into 1,684 train / 124 validation / 191 test, with zero shared documents. Target: a sourced short answer (median 18 words). Known issues: 7 substantive duplicate groups (18 records) plus one near-duplicate pair need a versioned fix, and none touches the development questions; the training data has no unanswerable examples; the 191 historical test records are exposed, so the final test uses fresh documents.
Decisions Made	October 8: (1) Selected Nemotron 3.5 Lightning, checkpoint 400, for the final test. Three adapters tied at 4/24 fully supported answers; Nemotron had complete outputs and the human-reviewed sample. (2) The final benchmark uses fresh official documents, with questions written by a local AI model and source-checked rather than by me. (3) Withdrew the first AI judge after it counted refusals as supported answers, and rescored with a corrected one.
Blockers & Risks Triggered	Unsupported or prohibited answer (fired): the adapter put an unsupported claim in 19 of 24 answers, and 0 of 30 reviewed citations were correct. Response: no promotion; on October 14 I changed the final test to supply the source passage. Train/test leakage (fired): the duplicate groups above; response: versioned fix, development results unaffected. Tool/export incompatibility (fired): Muse's output format broke the answer parser on 23 of 24 answers; response: Muse dropped. Timeline (fired): the final test slipped past Week 8; response: moved to Weeks 9–10, still before the October 25 deadline.
Next Period: Top Three Tasks	Before Milestone 2 (October 25): 1) Generate, source-check and lock the fresh benchmark, including unanswerable questions. 2) Run the one-time test of the Nemotron adapter against its base with the source passage supplied. 3) Rate the final outputs myself, run the robustness and subgroup analyses, and write Milestone 2.
Hours Logged	About 20 hours (my estimate; not timed).

AI Use Disclosure

Hermes Agent (GPT-6 Luna and GPT-6.1 Sol; Claude Opus 5.5 for the October 18 revision) drafted this log from my project records; local Qwen models produced the advisory development scores. I made the ratings, rulings and decisions reported here, estimated the hours, checked the log against the records and remain responsible for it.