AUtoMLy | Progress Log 2

INAI 6986 · Week 8: October 5–11, 2026 · J.O.

Model Evaluation week · Log due October 11; updated October 19 · Week 8 facts unchanged; later events carry their dates.

Log Number & Period	Log 2 · Week 8, October 5–11, 2026.
Planned vs. Completed	Planned (PM Plan, Week 8): evaluation metrics, validation strategy and error analysis for the trained models. Completed: (1) Week 8 error review: I rated 30 Nemotron outputs (15 validation questions, base and adapter; 180 ratings), October 6–7. (2) Rulings on all 16 cross-split duplicate-answer groups and one near-duplicate pair, October 5–6. (3) One prompt-refinement experiment on 8 questions, October 7; it failed and was not promoted. (4) A predeclared comparison of all four model families, base and adapter, on the same 24 validation questions (192 outputs), October 8. Not completed in Week 8: the one-time final test, moved to Weeks 9–10 because the fresh benchmark was not built yet; it ran October 19 (see Decisions Made and Risks).
Data-Readiness Status	1,999 answerable records; split by document into 1,684 train / 124 validation / 191 test, with zero shared documents. Target: a sourced short answer (median 18 words). Known issues: 7 substantive duplicate groups (18 records) plus one near-duplicate pair need a versioned fix, and none touches the development questions; the training data has no unanswerable examples; the 191 historical test records are exposed, so the final test uses fresh documents.
Decisions Made	October 8: (1) Selected Nemotron 3.5 Lightning, checkpoint 400, for the final test. Three adapters tied at 4/24 fully supported answers; Nemotron had complete outputs and the human-reviewed sample. (2) The final benchmark uses fresh official documents, with questions written by a local AI model and source-checked rather than by me. (3) Withdrew the first AI judge after it counted refusals as supported answers, and rescored with a corrected one. October 18: replaced the proposed Qwen final-test judge with OpenAI GPT Luna, paired with Claude Sonnet 5.5, so neither final judge shares Qwen's training-question lineage. Historical Qwen scores stay unchanged. October 19: ran the one-time final test on 150 fresh questions and did not promote the adapter (details below).
Blockers & Risks Triggered	Unsupported or prohibited answer (fired): the adapter put an unsupported claim in 19 of 24 answers, and 0 of 30 reviewed citations were correct. Response: no promotion; on October 14 I changed the final test to supply the source passage. Train/test leakage (fired): the duplicate groups above; response: versioned fix, development results unaffected. Tool/export incompatibility (fired): Muse's output format broke the answer parser on 23 of 24 answers; response: Muse dropped. Timeline (fired): the final test slipped past Week 8; response: moved to Weeks 9–10 and run October 19, before the October 25 deadline. October 19 final test: the adapter scored 111/120 on answerable questions against 114/120 for the untouched base (threshold: +10 points), declined 2 of 30 unanswerable questions (base 24; threshold 90%), and put an unsupported claim in 25 of 150 answers (threshold under 2%). Response: not promoted; the training data had no unanswerable examples.
Next Period: Top Three Tasks	Planned in Week 8 and finished October 19: the fresh benchmark, the one-time test and its scoring (two AI judges, my review, code citation checks, robustness and subgroup analyses). Next, before October 25: 1) Submit Milestone 2 with the final-test results. 2) Keep the final-test notebook and evidence files in the repository current. 3) Plan, but not start, a retrained adapter that includes unanswerable training examples.
Hours Logged	About 20 hours (my estimate; not timed).

AI Use Disclosure

Hermes Agent (GPT-6 Luna and GPT-6.1 Sol; Claude Opus 5.5, Claude Sonnet 5.5 and GPT-6.1 Sol for the October 18 revisions) drafted this log from my records. Local Qwen produced historical development scores. I made the ratings, rulings and decisions, estimated the hours, checked the log and remain responsible for it.